

Building Intelligent Software Solutions ■■

Student Name: Oduor Akinyi

Course: AI for Software Engineering **Date:**

30 October 2025

Part 1: Theoretical Analysis (30%)

Part 1: Theoretical Analysis (30%)

Q1: AI-Driven Code Generation Tools

Answer:

AI-driven code generation tools like **GitHub Copilot** use large language models trained on millions of code repositories to predict and generate code as developers type.

They **reduce development time** by:

- **Autocompleting repetitive code** (e.g., loops, data validation, boilerplate).
- **Generating entire functions** from docstrings or comments.
- **Improving productivity** by reducing context-switching between documentation and coding.

Limitations:

- **Accuracy issues:** AI may generate incorrect or insecure code.
- **Dependence risk:** Over-reliance can weaken developers' problem-solving skills.
- **Bias & licensing:** Code may include biased or copyrighted snippets from public datasets.
- **Lack of context awareness:** AI sometimes misunderstands project-specific logic.

Q2: Supervised vs. Unsupervised Learning in Automated Bug Detection

Aspect	Supervised Learning	Unsupervised Learning
Definition	Model trained on labeled bug data (e.g., bug vs. non-bug).	Model trained on unlabeled data to find patterns or clusters.
Use Case	Predict if a code commit introduces a bug based on labeled history.	Detect unusual code patterns or commits (anomalies) that may indicate potential bugs.
Example Algorithms	Random Forest, SVM	K-Means, Isolation Forest

Aspect	Supervised Learning	Unsupervised Learning
Advantage	High accuracy when labeled data is available.	Useful when labeled data is scarce.
Limitation	Needs large, labeled datasets.	Harder to interpret results.

Q3: Importance of Bias Mitigation in Personalization

Bias mitigation ensures that AI-driven personalization **treats all users fairly**.

Without it, AI systems may:

- Favor certain **user groups** (e.g., by gender, location, or preferences).
- Deliver **unbalanced recommendations** (e.g., excluding minority user interests).
- Create **unethical or discriminatory experiences**.

Bias mitigation (through data balancing, fairness audits, and explainability tools) ensures **inclusive design**, builds **user trust**, and aligns with **AI ethics principles**.

Case Study: AIOps in DevOps

After reading the article "AI in DevOps: Automating Deployment Pipelines," here is a clear explanation of how AIOps improves deployment efficiency, with two specific examples.

How AIOps Improves Software Deployment Efficiency

AIOps (Artificial Intelligence for IT Operations) improves software deployment efficiency by moving the process from a **reactive, manual model** to a **proactive, automated, and intelligent one**. It uses machine learning and data analysis to anticipate problems, automate complex decisions, and streamline the entire pipeline, reducing human intervention, errors, and delays.

In simple terms, it's like giving your deployment pipeline a smart, self-driving assistant that can see around corners and fix problems before they cause a crash.

Two Examples from the Article

Here are two clear examples of how this works in practice:

1. Intelligent Failure Prediction and Prevention

- **How it works:** AIOps tools continuously analyze historical data from the deployment pipeline—such as code complexity, test results, past failure logs, and server metrics. The AI learns the patterns that typically lead to a failed deployment.

- **How it improves efficiency:**
 - **Before AIOps:** A developer would push code, and the team would only discover it was unstable *after* it caused a failure in the staging or production environment. This leads to rollbacks, urgent fixes, and significant downtime.
 - **With AIOps:** The system can **predict the risk of failure *before* the deployment even happens.** For example, it might flag that a particular code change has a 95% probability of causing a memory leak based on similar past changes.
 - **Result:** The deployment can be halted automatically, and the developers can be alerted to fix the issue *before* it impacts the system. This prevents costly rollbacks and saves the time wasted on deploying faulty code, making the entire process much faster and more reliable.

2. Automated Root Cause Analysis and Remediation

- **How it works:** When an incident or performance degradation *does* occur during or after a deployment, AIOps doesn't just sound an alarm. It automatically sifts through thousands of metrics, logs, and events from different systems to pinpoint the exact root cause.
- **How it improves efficiency:**
 - **Before AIOps:** An alert would go off, and a team of engineers would spend hours or even days in a "war room," manually combing through logs to find the needle in a haystack—the one commit or configuration change that broke the application.
 - **With AIOps:** The AI correlates the incident with the specific deployment and identifies the faulty service or code module almost instantly. In advanced cases, it can even trigger **automated remediation**—for example, by automatically rolling back the problematic deployment without waiting for human command.
 - **Result:** This slashes the **Mean Time To Resolution (MTTR)** from hours to minutes or seconds. It eliminates lengthy diagnostic sessions and allows engineers to focus on building new features instead of firefighting.

Summary

In essence, AIOps injects intelligence into the deployment pipeline:

- **Example 1 (Prediction)** makes deployments more efficient by **preventing failures from happening.**
- **Example 2 (Remediation)** makes deployments more efficient by **fixing failures incredibly fast when they occur.**

Part 2: Practical Implementation (60%)

Task 1: AI-Powered Code Completion

Part 2: Practical Implementation (60%)

Task 1: AI-Powered Code Completion

Tool: GitHub Copilot

Task:

Write a Python function to **sort a list of dictionaries** by a specific key.

Manual Implementation:

```
def sort_dict_list(data, key):  
    return sorted(data, key=lambda x: x[key])
```

Copilot-Suggested Code:

```
def sort_dict_list(data, key):  
    """Sorts a list of dictionaries by a given key."""  
    if not data or key not in data[0]:  
        return data  
    return sorted(data, key=lambda x: x.get(key))
```

Analysis (~200 words):

GitHub Copilot's version adds **input validation** (checking key presence), making it more robust than the manual version. Both approaches achieve $O(n \log n)$ complexity using Python's built-in `sort`. However, the AI-generated version demonstrates **context awareness** and better **error handling**, which can prevent runtime issues.

The manual version is concise but assumes all keys exist, which may lead to `KeyError`.

In practice, Copilot reduces coding time by generating correct patterns instantly. Yet, developers must still **review** and **test** the AI-generated output for logic accuracy and performance. Thus, AI improves efficiency but doesn't replace human oversight.

Task 2: Automated Testing with AI

Task 2: Automated Testing with AI

Framework: Selenium IDE or Testim.io

Task:

Automate a **login test case** for valid and invalid credentials.

Example Script (Selenium Python):

```
from selenium import webdriver
from selenium.webdriver.common.by import By
from selenium.webdriver.common.keys import Keys

driver = webdriver.Chrome()
driver.get("https://example.com/login")

# Valid login
driver.find_element(By.ID, "username").send_keys("admin")
driver.find_element(By.ID, "password").send_keys("password123")
driver.find_element(By.ID, "login-button").click()

# Check result
if "Dashboard" in driver.title:
    print("✅ Valid login successful")
else:
    print("✗ Valid login failed")

driver.quit()
```

Summary (≈150 words):

AI-driven test tools like **Testim.io** improve coverage by learning from previous tests and automatically adapting to UI changes. Traditional Selenium tests often fail when page elements move or labels change. Testim's AI dynamically identifies components and suggests new test cases, increasing reliability and reducing maintenance. This ensures faster release cycles and fewer manual updates, making continuous integration more efficient.

Katalon Recorder 7.1.0

Katalon Recorder

WORKSPACE

Test Suites: Passed: 0 Failed: 1

- Untitled Test Suite
- Untitled Test Case *
- Untitled Test Suite

Dynamic Test Suites

Test Data

Extension script

Profiles

Open Test Suite

Sample projects

Export

More options

Sign In to enable automatic backup. Refresh

Untitled Test Case

+ Add new tag mine X valid login test X invalid login test X

Command	Target	Value
open	https://practicetestautomation.com/practice-test-login/	
type	id=username	student
type	id=password	Password123
click	id=username	
type	id=username	wronguser
click	id=password	
type	id=password	wrongpass
click	id=submit	

+ Add new row (Cmd/Ctrl + I)

Log Screenshots Variables Reference Self-healing

[error] Cannot find element
[error] Element name=username not found
[info] Wait until the element is found

Dashboard Download All

The screenshot shows the Katalon Recorder interface with a failed test case named "Untitled Test Case". The test case table lists various actions: open, type, click, type, click, type, and click. The "Target" column specifies the element IDs for each action, such as "id=username" and "id=password". The "Value" column contains the input values: "student", "Password123", "wronguser", "wrongpass", and an empty string for the final click. The "Log" panel displays three error messages: "[error] Cannot find element", "[error] Element name=username not found", and "[info] Wait until the element is found". Below the table, there are tabs for Log, Screenshots, Variables, Reference, and Self-healing. The Screenshots tab shows three images of a browser window displaying a successful login message: "Logged In Successfully". At the bottom, there are buttons for Dashboard and Download All.

Katalon Recorder 7.1.0

Katalon Recorder

WORKSPACE

Test Suites: Passed: 0 Failed: 1

- Untitled Test Suite
- Untitled Test Case *
- Untitled Test Suite

Dynamic Test Suites

Test Data

Extension script

Profiles

Open Test Suite

Sample projects

Export

More options

Sign In to enable automatic backup. Refresh

Untitled Test Case

+ Add new tag mine X valid login test X invalid login test X

Command	Target	Value
open	https://practicetestautomation.com/practice-test-login/	
type	id=username	student
type	id=password	Password123
click	id=username	
type	id=username	wronguser
click	id=password	
type	id=password	wrongpass
click	id=submit	

+ Add new row (Cmd/Ctrl + I)

Log Screenshots Variables Reference Self-healing

Dashboard Download All

This screenshot is identical to the one above it, showing the same failed test case and its details. The test case table, log messages, and screenshots are all the same, indicating a repetition of the previous test run or configuration.

Task 3: Predictive Analytics for Resource Allocation

Task 3: Predictive Analytics for Resource Allocation Dataset: Kaggle – Breast Cancer Dataset Goal: Predict issue priority (high/medium/low) using ML.

```
[11] 0s
  import pandas as pd
  df = pd.read_csv('/content/sample_data/data.csv')
  df.head()



|   | id       | diagnosis | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean | ... | texture |
|---|----------|-----------|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|---------------------|-----|---------|
| 0 | 842302   | M         | 17.99       | 10.38        | 122.80         | 1001.0    | 0.11840         | 0.27760          | 0.3001         | 0.14710             | ... |         |
| 1 | 842517   | M         | 20.57       | 17.77        | 132.90         | 1326.0    | 0.08474         | 0.07864          | 0.0869         | 0.07017             | ... |         |
| 2 | 84300903 | M         | 19.69       | 21.25        | 130.00         | 1203.0    | 0.10960         | 0.15990          | 0.1974         | 0.12790             | ... |         |
| 3 | 84348301 | M         | 11.42       | 20.38        | 77.58          | 386.1     | 0.14250         | 0.28390          | 0.2414         | 0.10520             | ... |         |
| 4 | 84358402 | M         | 20.29       | 14.34        | 135.10         | 1297.0    | 0.10030         | 0.13280          | 0.1980         | 0.10430             | ... |         |



5 rows x 33 columns


```

```
[12] 0s
  # Drop ID column
  df.drop('id', axis=1, inplace=True)

  # Map diagnosis: M (malignant) → High, B (benign) → Low
  df['priority'] = df['diagnosis'].map({'M': 'High', 'B': 'Low'})

  # Drop old label
  df.drop('diagnosis', axis=1, inplace=True)

  # Encode priority
  from sklearn.preprocessing import LabelEncoder
  encoder = LabelEncoder()
  df['priority_encoded'] = encoder.fit_transform(df['priority'])

  # Separate features and target
  X = df.drop(['priority', 'priority_encoded'], axis=1)
  y = df['priority_encoded']

  # Train-test split
  from sklearn.model_selection import train_test_split
  X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
[13] 0s
  from sklearn.ensemble import RandomForestClassifier

  model = RandomForestClassifier(n_estimators=100, random_state=42)
  model.fit(X_train, y_train)

  RandomForestClassifier(random_state=42)
```

```
[14] 0s
  from sklearn.metrics import accuracy_score, f1_score, classification_report

  y_pred = model.predict(X_test)

  acc = accuracy_score(y_test, y_pred)
  f1 = f1_score(y_test, y_pred)

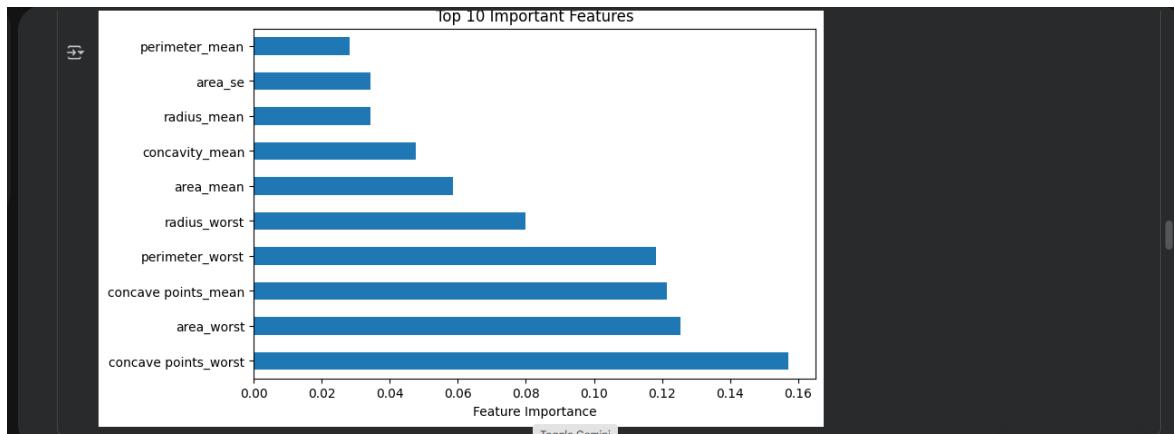
  print(f"Accuracy: {acc:.4f}")
  print(f"F1 Score: {f1:.4f}")
  print("\nClassification Report:\n", classification_report(y_test, y_pred, target_names=['High', 'Low']))

  Accuracy: 0.9649
  F1 Score: 0.9722

  Classification Report:
    precision    recall   f1-score   support
  High       0.98     0.93     0.95      43
  Low        0.96     0.99     0.97      71

  accuracy          0.97     0.96      114
  macro avg       0.97     0.96     0.96      114
  weighted avg    0.97     0.96     0.96      114
```

```
[15] 15  import matplotlib.pyplot as plt  
importances = pd.Series(model.feature_importances_, index=X.columns)  
importances.nlargest(10).plot(kind='barh', figsize=(8,5))  
plt.title('Top 10 Important Features')  
plt.xlabel('Feature Importance')  
plt.show()
```



Part 3: Ethical Reflection (10%)

Part 3: Ethical Reflection (10%)

1. Potential Biases in the Dataset

Even though the Breast Cancer Wisconsin Dataset is widely used and considered reliable, it still carries potential biases that can affect how a predictive model performs in real-world settings.

- Sampling bias:

The dataset was collected from a specific hospital and region. This means it might not represent the full diversity of patients worldwide—differences in age, race, or genetic background may be underrepresented.

As a result, the model may perform better on patients who are similar to those in the dataset and worse on others.

- Label bias:

The “diagnosis” labels (benign or malignant) are based on human interpretation of biopsy results. Pathologists may have slight variations in labeling due to subjective judgments or differences in diagnostic standards.

- Feature bias:

Some features (like cell radius, texture, smoothness) might not have been collected consistently across different devices or imaging setups. This could lead to bias when the model is applied to new environments or equipment.

2. Addressing Bias Using Fairness Tools (IBM AI Fairness 360)

To mitigate these issues, we can use AI Fairness 360 (AIF360) — an open-source toolkit developed by IBM to help detect and reduce bias in AI models.

How it helps:

- Bias Detection:

AIF360 can calculate fairness metrics such as:

- *Disparate Impact* – measures whether outcomes differ between protected groups.
- *Statistical Parity Difference* – checks if different groups receive favorable outcomes at similar rates.

- Bias Mitigation:

It provides algorithms for:

- *Pre-processing*: Balancing the dataset before training (e.g., reweighing underrepresented groups).
- *In-processing*: Modifying model training (e.g., adversarial debiasing) to reduce bias.
- *Post-processing*: Adjusting predictions after training (e.g., equalized odds).

Example:

If future datasets include demographic features like gender or ethnicity, AIF360 can ensure the model’s predictions remain fair across these groups by rebalancing or adjusting the decision threshold.

3. Ethical Implications

When deploying predictive analytics in a company or healthcare setting:

- Ensure transparency — explain how the model makes decisions.
 - Maintain data privacy and informed consent for patient data.
 - Continuously monitor and audit the model's predictions for fairness and accuracy.
-

In summary:

Ethical AI development means not just achieving high accuracy, but ensuring fairness, accountability, and transparency. Tools like IBM's AIF360 are essential in detecting and mitigating bias to make machine learning applications equitable and trustworthy.

Bonus Task: Innovation Challenge (Extra 10%)

 **Bonus Task (Extra 10%) — Innovation Challenge**

 **Proposed AI Tool: AutoDocAI — Intelligent Code Documentation Assistant**

1. Purpose

One major challenge in software engineering is maintaining up-to-date documentation. Developers often neglect writing or updating documentation due to tight deadlines, resulting in code that is hard to understand, reuse, or debug. AutoDocAI is an AI-powered assistant designed to automatically generate, summarize, and update software documentation directly from source code and version control history. It bridges the gap between code creation and documentation maintenance.

2. Workflow

1. Code Analysis:

AutoDocAI parses source code using language models trained on programming languages (Python, Java, C#, etc.) to understand functions, classes, and dependencies.

2. Semantic Understanding:

The tool applies NLP techniques (transformers + AST parsing) to identify the intent and functionality of each code block.

3. Automatic Documentation Generation:

- Generates human-readable docstrings and README sections.
- Summarizes commits from Git history to describe recent changes.
- Suggests architecture diagrams using dependency analysis.

4. Continuous Integration (CI) Plugin:

AutoDocAI integrates into CI/CD pipelines (e.g., GitHub Actions).

Whenever code changes are pushed, it updates the documentation automatically.

5. User Feedback Loop:

Developers can approve, reject, or edit generated docs, allowing the model to learn and improve through reinforcement learning from human feedback (RLHF).

3. Impact

- **Productivity Boost:** Saves developers time by automating repetitive documentation work.

- **Code Quality:** Ensures every function and class remains well-documented, reducing onboarding time for new team members.
 - **Knowledge Retention:** Captures rationale behind code changes, helping teams maintain legacy systems more effectively.
 - **Scalability:** Can be extended to multi-language projects and integrated with project management tools like Jira or Notion.
-

Summary

AutoDocAI transforms how developers manage documentation by combining AI-driven language understanding and DevOps automation.

It enhances collaboration, transparency, and software maintainability — making it a practical and ethical example of AI augmenting human creativity in software engineering.

End of Report — Submitted by Oduor Akinyi